

Candidate questions (5 strategies × 5 generators)

Team A — Static integrity

A1 LexicalHygiene

A2 BiasStats

A3 FactEcho

A4 TemplateFingerprint

Team B — Tri-judge LLM panel

B1 TriJudgeAnswer

B2 ClosedBookSolvability

B3 UbiquityRisk

B4 Ambiguity (esc)

B5 VerifierSkip (esc)

Team C — Deterministic checks

C1 DistractorDifficulty (esc)

C2 CategoryLeak

C3 SourceSwap (esc)

C4 DifficultyAudit

Team D — Corpus statistics

D1 SelfPreference

D2 DedupCalibration (esc)

D3 SkewAudit

Each agent calibrated against human gold sheet via Cohen's κ — signals with $\kappa < 0.6$ are advisory-only

DROP (critical FAIL) → 341 questions removed

RELABEL (C4 / human) → 1,259

KEEP → release_v1.2