
OenoBench: A Wine-Domain Benchmark for Knowledge-Grounded Evaluation of Large Language Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We introduce **OenoBench**, a wine-domain knowledge benchmark of **3,266**
2 **multiple-choice questions** across six pillars (regions, grape varieties, viticulture,
3 winemaking, producers, business) and four difficulty tiers. The corpus is built from
4 **38,104 atomic, source-anchored facts** extracted by 35 provenance-verified scrap-
5 ers from government registries (INAO, TTB, OIV), peer-reviewed journals, and
6 Wikipedia/Wikidata. Our methodological contribution is an *LLM-driven pipeline* in
7 which language models reformat verified facts and audit the result, but never serve
8 as the source of truth: every claim traces to a URL, every question is generated by
9 one of five strategies across five generator families, and every question is scored by
10 a nine-agent audit calibrated against a human gold sheet via Cohen’s κ . Evaluating
11 sixteen frontier configurations, we find: (i) overall accuracy spans 53%–84%, led
12 by **o3 at 83.6%**; (ii) reasoning-mode lift concentrates in **DeepSeek R1 (+6.8 pp)**
13 and is absent in Claude Opus and Gemini Pro; (iii) Anthropic shows **+9 pp self-**
14 **preference** on its own questions while Google shows **−8 pp** inverse preference;
15 (iv) frontier open-weight models share the cost-vs-accuracy Pareto frontier with
16 proprietary reasoning models; and (v) every config gains **≈33 pp** on closed-book
17 solvable items, revealing a parametric-recall ceiling that only the contextual slice
18 avoids. We release corpus, audit findings, human-review app, and construction
19 code under CC-BY-SA-4.0.

20 1 Introduction

21 Large language models are increasingly evaluated on knowledge-intensive benchmarks, but the most
22 widely-used resources — MMLU [6], GPQA [16], and their successors [19, 21, 15] — share three
23 weaknesses: they aggregate breadth at the expense of *depth* within any single discipline; they are
24 typically authored by a small team using one generation pipeline, exposing systematic blind-spots
25 and stylistic regularities that frontier models can learn to exploit [14, 20]; and the questions are
26 increasingly suspected of contamination through web-scale pre-training corpora [17, 13, 3]. As
27 models grow more capable, the marginal value of a benchmark depends increasingly on *how* it was
28 constructed.

29 We argue that domain-specialised benchmarks built around bodies of knowledge that are *externally*
30 *validated by humans* — through certification, regulation, or peer-reviewed practice — offer a useful
31 complement to broad knowledge benchmarks. Examples already exist in medicine [9, 10], law [5],
32 and finance [8]. We extend this line of work to a domain that has so far received little attention from
33 the benchmarking community, despite being unusually well-defined: **wine**.

Why wine. Wine integrates plant biology, geology and climate science, organic chemistry, food microbiology, sensory neuroscience, regulatory law, and business economics. It is one of very few domains for which inter-governmental and governmental sources publish structured taxonomic data (INAO, OIV, TTB, regional consortia), university research groups (UC Davis, Geisenheim, Bordeaux Sciences Agro) produce open scholarly literature, and professional certification bodies (WSET Levels 1–4, Court of Master Sommeliers, Institute of Masters of Wine) define a graded competence ladder. These properties make it possible to build a benchmark whose every question is traceable to a verifiable source while still spanning the recall–reasoning–application spectrum that distinguishes a strong taster, winemaker, or buyer from a memorised glossary.

Contributions.

1. **A 3,266-question wine benchmark** across six pillars and four difficulty tiers, built from 38,104 atomic provenance-tagged facts (Section 3).
2. **An LLM-driven, fact-grounded generation pipeline** with five strategies $\times$ five generator models, per-strategy quotas, and an explicit closed-book pre-screen — LLMs reformat verified facts, never supply them.
3. **A nine-agent automated audit** (Section 4) calibrated against a human gold sheet via Cohen’s κ ; the audit removed 341 questions and re-labelled difficulty on 1,259, and surfaced one agent (closed-book solvability) that over-reports leakage relative to humans, which we keep with explicit disclosure rather than drop.
4. **A bias-aware evaluation across sixteen frontier configurations** (Section 5) reporting overall accuracy, reasoning-mode lift, a Self-Preference Score, cost-efficiency Pareto, and a closed-book vs. source-grounded contrast that isolates parametric recall from contextual reasoning.

OenoBench — corpus, audit findings, construction code, and human-review app — ships under CC-BY-SA-4.0; release URLs, reproducibility scripts, and the per-source licensing table are collected in Appendix E.

2 Related Work

Broad knowledge benchmarks. MMLU [6], BIG-Bench [19], AGIEval [21], GPQA [16], HELM [12], and Humanity’s Last Exam [15] measure factual breadth and reasoning across many subjects. We instead maximise depth within a single externally-validated knowledge domain and treat *construction-time bias control* as a first-class contribution rather than a downstream decontamination step.

Domain-specific benchmarks. MedQA [9], PubMedQA [10], LegalBench [5], and FinanceBench [8] ground evaluation in regulated professions with clear sources of truth. We adopt their expert-syllabus / source-traceable backbone and add multi-model generation, multi-agent audit, and self-preference reporting. OenoBench is, to our knowledge, the first such benchmark for wine.

LLM-as-judge bias and self-preference. A growing literature documents that models tend to favour outputs stylistically similar to their own [14, 20, 18], biasing benchmarks where the same model generates and grades. We mitigate by distributing generation across five generator families plus deterministic templates and by reporting per-model self-vs-other accuracy gaps as a first-class diagnostic (Section 5.3).

Benchmark contamination and solvability. Pre-training exposure inflates apparent capability [17, 13, 3]. We mitigate by generating from atomic facts (a fact-echo agent flags any question whose longest common substring with the source exceeds 65%) and by running an explicit *closed-book solvability* pre-screen and audit (Section 5.5). The same audit, calibrated against a human gold sheet, reveals that LLM judges over-report leakage by an order of magnitude — a finding with broader implications for the LLM-as-judge paradigm.

Wine in NLP. Wine has previously appeared in NLP via review-text recommendation [2], taxonomy learning [11], and judge-reliability studies [7]. OenoBench is the first benchmark targeting *structured factual knowledge* of viticulture, winemaking, and wine regions evaluated against frontier LLMs.

3 Dataset Construction

OenoBench is constructed by an LLM-driven pipeline (Figure 1) in two stages: (1) a *fact-collection* stage that builds a corpus of atomic, provenance-tagged statements from authoritative sources, and (2) a *question-generation* stage that converts those facts into benchmark questions using five complementary strategies and five generator models. The pipeline is designed around a single methodological commitment: large language models are used as *rephrasers and auditors*, never as the source of truth. Every fact in the corpus traces to a URL with a tier-of-authority label, every question is anchored in one or more such facts, and every question is scored by a nine-agent automated audit (Section 4) before release.

3.1 Domain taxonomy




Figure 1: End-to-end OenoBench pipeline: authoritative sources → 35 scrapers → atomic facts → closed-book gate → nine-agent audit → release.

We organise wine knowledge into six *domain pillars*, chosen to align with the WSET Level 3 and Diploma syllabi: (i) **wine regions** (appellations, classifications, geography, climate, soils); (ii) **grape varieties** (ampelography, parentage, regional distribution, sensory profile); (iii) **viticulture** (vineyard practice, training systems, pests and diseases, vintage effects); (iv) **winemaking** (vinification choices, fermentation, ageing, faults); (v) **producers** (estates, négociants, co-operatives, key houses); and (vi) **wine business** (markets, regulation, distribution, economics). Each domain has a separate target share informed by syllabus weighting and source availability; the realised distribution is reported in Table 1.

3.2 Source tiering, licensing, and non-fabrication

Every fact is tagged with one of three source tiers, mirroring evidence-based-medicine practice:

- **Tier 1 (official, 19.6%)**: government registries and inter-governmental bodies (INAO, OIV, TTB, regional consortia, the UC Davis AVA database).
- **Tier 2 (authoritative, 76.6%)**: Wikipedia, Wikidata SPARQL, wine-body and consortium websites, peer-reviewed journals (*OENO One*, *Vitis*, *Australian Journal of Grape and Wine Research*).
- **Tier 3 (reliable, 3.8%)**: curated open datasets on HuggingFace and Kaggle, secondary reference databases.

Licensing and non-fabrication. Every source is public-record or released under a licence that permits scraping and redistribution: Tier-1 government data under explicit open-data licences (INAO Licence Ouverte, OIV, UC Davis AVA, TTB public-domain), Tier-2 Wikipedia/Wikidata under CC-BY-SA / CC0 and journals (*OENO One*, *Vitis*, *AJEV*) open-access; the full per-source licence

table is Table 16 in Appendix E. A central design commitment is that *no fact in the corpus comes from an LLM’s internal knowledge*: every fact is extracted by a source-specific scraper, stamped with source URL, retrieval timestamp, and tier label before insertion, and the fact-insertion utility rejects records without a verified source-URL trace. We emphasise this because LLM-generated content masquerading as scraped data is a contamination vector that conventional dataset documentation does not detect.

Table 1: Atomic-fact corpus: 38,104 facts from 35 scrapers, six domain pillars, three source tiers.

Domain pillar	Facts	Share	Tier 1	Tier 2/3
Wine regions	18,943	49.7%	23.4%	76.6%
Producers	6,215	16.3%	9.1%	90.9%
Grape varieties	5,959	15.6%	18.8%	81.2%
Viticulture	3,635	9.5%	12.0%	88.0%
Wine business	1,985	5.2%	31.4%	68.6%
Winemaking	1,367	3.6%	16.7%	83.3%
Total	38,104	100.0%	19.6%	80.4%

3.3 Atomic fact extraction

Source pages are decomposed into atomic facts by a five-stage pipeline (Figure 5, Appendix B.1): sentence splitting; reference resolution (pronouns replaced by their entity referents); domain classification into one of the six pillars; a length/predicate validator (5–50 words, must contain a verb, no dangling references); and a region-keyword on-topic filter that prevents cross-contamination (e.g. Austrian content in a Bordeaux scraper). Wikidata SPARQL queries use direct country relations (P17) rather than transitive parents (P131*), the latter having caused severe cross-region leakage in early scrapers (Appendix B).

Facts are required to be *atomic* (one assertion per sentence), *entity-tagged* (subject and object linked to canonical knowledge-graph IDs), and *source-faithful* (a paraphrase, never a verbatim copy); the last is enforced at extraction and re-checked at audit time by agent A3.

3.4 Question generation: five strategies, five models

A single-strategy, single-model benchmark inherits the blind-spots of its pipeline. We mitigate this by combining *five strategies* and *five generator models*, with per-strategy and per-model quotas enforced by an orchestrator that resumes safely across runs. Strategies are:

- **Fact-to-question (58.4%, 1,909 questions)**: the LLM rewrites a single verified fact as a multiple-choice question; preserves grounding but tends toward recall.
- **Distractor mining (12.4%, 405)**: confusable-entity sampling produces wrong answers that share the right answer’s category and dimension, raising plausibility and forcing fine-grained discrimination.
- **Template (11.9%, 389)**: forty-five deterministic parameterised templates across the six pillars; pure entity substitution, zero LLM creativity — a baseline against which neural generation is measured.
- **Scenario synthesis (9.8%, 319)**: coherent fact clusters are converted into applied multi-fact reasoning prompts (e.g. a sommelier service decision, a winemaker blending decision, a viticulturist vintage call). Domain-specific scenario types prevent persona-content mismatch.
- **Comparative (7.5%, 244)**: entity-affinity scoring pairs related entities (e.g. two Burgundy premiers crus, two Rhône grapes) and asks “which differs in X” or “what do both share”.

The five generator models are Claude Opus 4.7, GPT-5, Gemini 2.5 Pro, Llama 3.1 405B, and Qwen 3.5 235B, each accessed through a unified OpenRouter client with consistent temperature and top-*p*. The realised generator distribution across the released corpus is **Qwen 20.4%, Llama 19.3%, Claude 19.0%, ChatGPT 16.6%, Gemini 12.9%, deterministic templates 11.9%** (computed directly from the 3,266 released questions). The orchestrator allocates generation share dynamically based on per-pilot audit pass rates — generators whose questions pass audit reliably get a larger share in

subsequent rounds — subject to a hard **per-generator cap of 21 %** to keep any single model from dominating the corpus.

3.5 Final corpus: release_v1.2

After generation, audit, and post-eval review, the released corpus release_v1.2 contains **3,266 questions**. The distribution along the domain and difficulty axes is:

Table 2: Composition of the released corpus.

Domain	N	%	Difficulty	N	%
Wine regions	1,108	33.9%	L1 (entry)	694	21.2%
Grape varieties	766	23.5%	L2 (intermediate)	894	27.4%
Producers	515	15.8%	L3 (advanced)	678	20.8%
Viticulture	502	15.4%	L4 (expert)	1,001	30.6%
Wine business	250	7.7%			
Winemaking	187	5.7%			
Total	3,266	100%	Total	3,266	100%

release_v1.2 ships under CC-BY-SA-4.0 on HuggingFace with a Croissant manifest [1] and a full Datasheet [4] in Appendix A.

4 Multi-Agent Quality Audit

The audit pipeline runs every generated question through nine automated agents in four teams. Each agent emits a per-question {PASS, WARN, FAIL} signal calibrated against a human-reviewed gold sheet via Cohen’s κ ; agents whose agreement falls below threshold are downweighted to advisory-only. The released human-review web app (Appendix D) collects the gold-sheet ratings, and the team-by-team architecture (Figure 6, Appendix C) traces the flow from candidate question to drop / relabel / keep verdict.

4.1 Team architecture

Of the 14 agents (Tables 3–4), **nine are always-run** on every candidate question; the remaining **five are escalation-gated** and are invoked only when an upstream agent flags a corpus-level concern. This deferred-execution policy keeps routine audit cost bounded ($\sim$ \$76 for the full corpus audit reported here, 5h 22m wall) while preserving full coverage when a defect class is suspected. The full audit code, prompts, and thresholds are released alongside the corpus (Appendix B).

Table 3: Audit agents (1/2): static and tri-judge teams. ALWAYS: run on every question. ESC: escalation-gated.

Team	Agent	Run	Function
A. Static	A1 LexicalHygiene	ALWAYS	Regex sweep for vague phrasing, marketing language, meta-questions.
	A2 BiasStats	ALWAYS	χ^2 on correct-answer position; Mann–Whitney U on correct-vs-distractor length.
	A3 FactEcho	ALWAYS	Longest-common-substring ratio between question text and source fact (FAIL at $LCS \geq 0.65$).
	A4 TemplateFingerprint	ALWAYS	POS-bigram logistic regression detector for template-induced stylistic regularities (held-out AUC 0.84).
B. Tri-judge	B1 TriJudgeAnswer	ALWAYS	Claude / GPT / Gemini panel reads question + source fact and votes on the marked-correct option.
	B2 ClosedBookSolvability	ALWAYS	Same panel, source fact <i>withheld</i> : can the question be answered from world knowledge alone?
	B3 UbiquityRisk	ALWAYS	Internationally-grown grape stem $\times$ region-class answer $\Rightarrow$ ambiguity flag.
	B4 Ambiguity	ESC	Tri-judge re-read flags questions with >1 defensible answer; invoked when B1 dissent rate elevated.
	B5 VerifierSkip	ESC	Self-consistency probe across re-prompted panels; invoked on disputed B1 verdicts.

Table 4: Audit agents (2/2): deterministic and corpus-statistics teams.

Team	Agent	Run	Function
C. Determin.	C1 DistractorDifficulty	ESC	Embedding-distance distractor analysis; invoked when A2 reports length skew.
	C2 CategoryLeak	ALWAYS	Wine-type distractor validation; flags red/white/sparkling category mismatches.
	C3 SourceSwap	ESC	Substitutes alternate source facts to test answer stability; invoked on B1/B2 disagreement.
	C4 DifficultyAudit	ALWAYS	Gemini Pro re-rates difficulty; FAIL/WARN when delta ≥ 2 from assigned.
D. Corpus stats	D1 SelfPreference	ALWAYS	Per-generator-family own-vs-other accuracy on a held-out probe; corpus-level signal.
	D2 DedupCalibration	ESC	Near-duplicate audit at cosine ≥ 0.92 ; invoked on suspicious cluster signals.
	D3 SkewAudit	ALWAYS	Country / sub-domain over-representation χ^2 .

4.2 Gold-sheet calibration

A gold sheet of 50 stratified questions per release is rated independently by three WSET-certified reviewers (Diploma, Level 3, Level 2 — one of whom is among the authors) through the human-review web app (Appendix D); **15+ rounds** (audit pilots v1–v16) form a calibration cycle covering eight rubrics (answer-correctness, distractor plausibility, ambiguity, source-faithfulness, needs-source, vague language, label correctness, verbatim copying). For each agent we compute Cohen’s κ against the rubric; **signals with $\kappa < 0.6$ are downweighted to advisory-only**.

The most consequential result is agent **B2 ClosedBookSolvability**: the human solved $\approx 12\%$ of gold-sheet questions without the source fact while the LLM tri-judge panel reported $\approx 83\%$ ($\kappa \approx 0.007$). We read this as a property of the *evaluator* — frontier LLM judges have absorbed enough wine knowledge during pre-training to over-attribute “world knowledge” to non-trivial questions. We retain the 1,601 B2-flagged questions with disclosure and turn the blind spot into a calibrated memorisation-reliance diagnostic via the closed-book vs. source-grounded contrast (Section 5.5).

4.3 Release-cycle results

The 9-agent audit on `release_v1.1` (3,670 candidate questions) produced the verdicts summarised in Table 6 (Appendix B.7): A1, A3, B1, C2 and B3 between them flagged 341 distinct questions for DROP, while C4 re-labelled difficulty on 1,259 questions and B2 flagged 1,601 as closed-book solvable but kept them with disclosure.

After drops, a follow-up post-evaluation pass (Section 5 and Appendix C) audited the 97 questions that all 16 evaluation configs answered incorrectly: 54 were removed as defects (wrong ground truth, all-correct options, duplicate options), 9 were dropped on borderline-review and the remaining 34 retained as legitimately hard. The final released corpus is 3,266 questions.

The C4-driven difficulty re-label was decisive: it shifted the corpus from 14% L3+L4 items (the original generator-assigned distribution) to **51% L3+L4** (Table 7, Appendix B.7), bringing the hardest-tier counts above the entry-tier counts and substantially sharpening difficulty-stratified analyses (Section 5.1).

5 Evaluation

We evaluate **16 frontier configurations** on the 3,266-question release in a single end-to-end run. The slate combines six within-family cost pairs (Claude Opus 4.7 vs. Claude Haiku 4.5, Gemini 2.5 Pro vs. Gemini 2.5 Flash, GPT-5 vs. GPT-5-mini, Llama 3.3 70B vs. Llama 3.1 8B, Qwen 2.5 72B vs. Qwen 2.5 7B), four reasoning-mode configurations (o3, DeepSeek R1, Gemini 2.5 Pro thinking, Claude Opus 4.7 thinking), and two additional standard models (DeepSeek V3, Mistral Large). Every config sees the same 3,266 questions, with single-letter output (A–D) at `max_tokens=5` and a five-stop fallback. We organise the evaluation around five questions, one per subsection.




Figure 2: Overall accuracy on the 3,266-question corpus across 16 configurations. Reasoning-mode configs are hatched.

5.1 Overall capability ranking

OenoBench cleanly stratifies the slate into three bands (Figure 2): **frontier** ($\geq 80\%$, led by **o3** at **83.6%**, with GPT-5, Gemini 2.5 Pro, and Claude Opus 4.7 clustered within 3 pp); **mid-tier** (mid-60s to high-70s); and **small-tier** (low-50s to low-60s). No config performs at random; the best-vs-worst spread is 30 pp, comparable to GPQA [16] but on a domain corpus $30\times$ smaller and fully provenance-tagged. Difficulty-stratified accuracies (Appendix C, Table 14) confirm a clean L1→L4 gradient: 93.6% on L1 down to **58.7% on L4 expert items**, with frontier configs holding 68–71% on L4 and small configs at 39–45%.

5.2 Reasoning-mode lift

Reasoning-mode lift is concentrated in a single family (Figure 7, Appendix C). Only **DeepSeek R1 vs. V3** (+6.8 pp [+4.6, +8.8]) has a CI excluding zero; the three frontier pairs (o3 vs. GPT-5, Gemini Pro thinking vs. Pro, Claude Opus thinking vs. Opus) are all statistically indistinguishable from zero. We attribute the asymmetry to headroom: V3 leaves substantial parametric ceiling at 70.3% which chain-of-thought can recover, while frontier models are already near their recallable ceiling. Per-difficulty breakdowns (Appendix C, Table 8) show the frontier reasoning configs gaining at most 1–2 pp at L4 and flat elsewhere, making the cost-vs-lift trade unfavourable except for DeepSeek (Section 5.4).

5.3 Self-preference bias

We compute the *Self-Preference Score (SPS)* per config m as $SPS(m) = \text{Acc}(m \mid Q_{\text{own}}) - \text{Acc}(m \mid Q_{\text{other}})$ where Q_{own} collects questions generated by m ’s family and Q_{other} its complement. A pipeline that successfully neutralises generator fingerprints should produce $|SPS|$ near zero.

The result (Figure 3) is family-dependent: **Anthropic** clusters at +9 to +10 pp (all CIs above zero), **OpenAI** is statistically zero, and **Google** shows an unexpected *negative* −6 to −10 pp cluster. The cross-family matrix (Figure 3, left) traces the asymmetry: Google-generated questions are uniformly harder for everyone, while Anthropic-generated questions are uniformly easier for Anthropic only. We interpret the Anthropic gap as a residual stylistic fingerprint that pre-release paraphrase passes did not neutralise, and the Google gap as harder-by-phrasing rather than by leak; both are first-class disclosures with the released corpus.




Figure 3: Self-preference bias. *Left*: cross-family accuracy matrix (rows = evaluator, columns = generator). *Right*: per-config SPS $\delta = \text{Acc}(\text{own}) - \text{Acc}(\text{other})$ with 95% bootstrap CI.




Figure 4: Total OpenRouter cost (log scale) versus accuracy on the full 3,266-question slate. Family colours match Figure 2; the Pareto frontier (red dashed) illustrates the cost-vs-quality trade-off.

236 5.4 Cost efficiency

237 The slate spans nearly four orders of magnitude in evaluation cost (Figure 4). Five configs sit on
 238 the Pareto frontier spanning the cost-quality spectrum: **Llama 3.1 8B** (60.5 % / \$0.01), **Gemini**
 239 **2.5 Flash** (75.1 % / \$0.12), **GPT-5-mini** (78.4 % / \$2.82), **Claude Opus 4.7** (81.0 % / \$3.35), and
 240 **o3** (83.6 % / \$11.80). The two most expensive configs (Gemini 2.5 Pro \$29.47, GPT-5 \$21.90)
 241 are dominated by cheaper Pareto neighbours: their extra spend buys output tokens, not better wine
 242 answers. Reasoning-mode lift (Section 5.2) yields a Pareto win only for DeepSeek R1, which is itself
 243 dominated by GPT-5-mini. The efficient boundary therefore favours small proprietary plus a single
 244 reasoning model over reasoning-everywhere or scale-everywhere strategies.

245 5.5 Closed-book vs. source-grounded performance

246 The B2 agent reported ~83% closed-book solvability against the human reviewer’s ~12% (Sec-
 247 tion 4.2) — itself a measurement of what frontier LLMs already *know* about wine. We exploit it by
 248 partitioning the corpus into 1,601 B2-flagged and 1,665 contextual / source-grounded questions and
 249 reporting per-config accuracy on each slice (Figure 8, Appendix C.3).

250 The contrast is the largest single effect we observe: the mean per-config closed-book minus contextual
 251 gap is +**32.6 pp** (range +26.6 to +39.6 pp, every CI above zero). The B2-flagged slice overlaps

heavily with parametric wine knowledge acquired in pre-training; the contextual slice is where wine knowledge must be reasoned through the source fact — **Claude Opus 4.7 and o3 hold $\sim 70\%$ on the contextual slice, small models drop to $\approx 45\%$** , and the inter-config spread widens from 27 to 33 pp. The contextual slice is the more discriminative test of wine reasoning; the B2-flagged slice is a calibrated indicator of how much apparent “wine knowledge” is parametric recall.

6 Applications and Future Work

6.1 Closing the wine-knowledge gap in low-cost models

Small open-weight models trail proprietary frontier ones by 20–30 pp on this benchmark, while the broad-benchmark gap is only 10–15 pp (Section 5.4). The 38,104-fact corpus is a natural supervised-fine-tuning target via a two-stage recipe: instruction tuning on held-out question/source-fact pairs, then LoRA adaptation on the full fact-question-source triples in the release schema. Both stages run on consumer GPUs without proprietary weights or pre-training compute. The contextual slice (Section 5.5) is the harder-to-game evaluation surface, so measured gain tracks genuine wine reasoning. A Llama-3.3-70B-class adaptation could plausibly close most of the GPT-5-mini gap at $\approx 1\%$ of proprietary inference cost.

6.2 Case-based evaluation with industry partners

Multiple-choice tests have ceilings: a perfect OenoBench score does not establish useful behaviour in real workflows. We are scoping an **applied case-based dataset** with industry partners along three personas:

- *Sommelier* $\rightarrow$ *customer service*: pairing prompts rated on recommendation, justification, and substitutions.
- *Winemaker* $\rightarrow$ *blending decisions*: vintage write-ups, lab analyses, and regulatory constraints in context; rated on technical correctness, compliance, and stylistic intent.
- *Viticulturist* $\rightarrow$ *vineyard decisions*: weather, soil/disease, canopy/yield records; recommendations rated against decisions actually taken in the same vintage.

6.3 Generalising the scrape-to-audit pipeline

The pipeline modules — tier-of-authority taxonomy, atomic-fact extraction, multi-strategy multi-model generation with per-model quotas, κ -calibrated multi-agent audit, and closed-book pre-screen — are domain-agnostic. Three near-term targets share wine’s structural features: *horticulture/agronomy* (regulated appellations, peer-reviewed journals, certification ladders); *regulated medicine* (clinical guidelines, drug formularies; current benchmarks rely on board-exam questions misaligned with practice); and *financial regulation* (FASB, IFRS, jurisdictional filings; CFA/CPA ladders). The released code reduces months of expert authoring to weeks of curation.

7 Limitations

Snapshot vs. moving target. Wine regulation, classifications, and producer ownership change on a multi-year cadence (INAO appellation revisions, DOCG promotions, estate sales). The released corpus is a snapshot dated 2026-04-01 and will require periodic re-extraction; we release the scrapers and provenance metadata so this is mechanically reproducible.

Closed-book audit calibration. The B2 ClosedBookSolvability agent shows $\kappa \approx 0.007$ with humans on the gold sheet (Section 4.2). Rather than drop the affected questions we keep them with explicit disclosure and report the closed-book vs. contextual contrast (Section 5.5) as a calibrated memorisation-reliance diagnostic.

Self-preference is not fully neutralised. The Anthropic +9 pp positive SPS (Section 5.3) shows the multi-model strategy did not fully neutralise stylistic fingerprints; downstream Anthropic-vs-other comparisons on the full corpus should be read with the per-config SPS in mind.

References

- [1] Mubashara Akhtar, Omar Benjelloun, Costanza Conforti, et al. Croissant: A metadata format for ML-ready datasets. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2024.
- [2] Li Chen, Yi-Ching Chen, Yiqun Liu, and Shaoping Ma. Recommending wines based on subjective user ratings. In *Proceedings of the 8th ACM Conference on Recommender Systems (RecSys)*, 2014.
- [3] Jesse Dodge, Maarten Sap, Ana Marasović, William Agnew, Gabriel Ilharco, Dirk Groeneveld, Margaret Mitchell, and Matt Gardner. Documenting large webtext corpora: A case study on the Colossal Clean Crawled Corpus. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2021.
- [4] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12), 2021.
- [5] Neel Guha, Julian Nyarko, Daniel E. Ho, Christopher Ré, Adam Chilton, et al. LegalBench: A collaboratively built benchmark for measuring legal reasoning in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [6] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations (ICLR)*, 2021.
- [7] Robert T. Hodgson. An examination of judge reliability at a major u.s. wine competition. In *Journal of Wine Economics*, volume 3, 2008.
- [8] Pranab Islam, Anand Kannappan, Douwe Kiela, Rebecca Qian, Nino Scherrer, and Bertie Vidgen. FinanceBench: A new benchmark for financial question answering, 2023.
- [9] Di Jin, Eileen Pan, Nassim Oufattole, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. What disease does this patient have? a large-scale open domain question answering dataset from medical exams. *Applied Sciences*, 11(14), 2021.
- [10] Qiao Jin, Bhuwan Dhingra, Zhengping Liu, William Cohen, and Xinghua Lu. PubMedQA: A dataset for biomedical research question answering. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- [11] Els Lefever, Iris Hendrickx, Ilja Croijmans, Asifa Majid, and Antal van den Bosch. A hybrid approach to domain-independent taxonomy learning applied to the wine domain. In *Language Resources and Evaluation (LREC)*, 2018.
- [12] Percy Liang, Rishi Bommasani, Tony Lee, et al. Holistic evaluation of language models. In *Transactions on Machine Learning Research (TMLR)*, 2023.
- [13] Inbal Magar and Roy Schwartz. Data contamination: From memorization to exploitation. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2022.
- [14] Arjun Panickssery, Samuel R. Bowman, and Shi Feng. LLM evaluators recognize and favor their own generations. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [15] Long Phan, Alice Gatti, Ziwen Han, Nathaniel Li, Josephina Hu, et al. Humanity’s last exam, 2025.
- [16] David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R. Bowman. GPQA: A graduate-level google-proof q&a benchmark. In *Conference on Language Modeling (COLM)*, 2024.
- [17] Oscar Sainz, Jon Ander Campos, Iker García-Ferrero, Julen Etxaniz, Oier Lopez de Lacalle, and Eneko Agirre. NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark. In *Findings of the Association for Computational Linguistics: EMNLP*, 2023.

- [18] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Aspell, Samuel R. Bowman, et al. Towards understanding sycophancy in language models. In *International Conference on Learning Representations (ICLR)*, 2024.
- [19] Aarohi Srivastava et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *Transactions on Machine Learning Research (TMLR)*, 2023.
- [20] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [21] Wanjun Zhong, Ruixiang Cui, Yiduo Guo, Yaobo Liang, Shuai Lu, Yanlin Wang, Amin Saied, Weizhu Chen, and Nan Duan. AGIEval: A human-centric benchmark for evaluating foundation models. In *Findings of the Association for Computational Linguistics: NAACL*, 2024.

A Datasheet for OenoBench

We follow the datasheet template of Gebru et al. [4].

A.1 Motivation

For what purpose was the dataset created? To evaluate factual knowledge of large language models on a multi-disciplinary, externally-validated domain (wine), and to provide a benchmark whose construction includes explicit bias controls (multi-model generation, multi-agent audit, self-preference scoring, closed-book vs. source-grounded partition).

Who created the dataset? The dataset was created by the anonymous author(s) for the NeurIPS 2026 Evaluations & Datasets Track. One of the authors holds the WSET Diploma in Wines, the highest pre-Master-of-Wine qualification of the Wine & Spirit Education Trust. Gold-sheet ratings used to calibrate the nine audit agents were produced by **three WSET-certified reviewers** (Diploma, Level 3, Level 2), one of whom is among the authors. Author identities are withheld for double-blind review.

Funding. The work was independently funded by the authors’ affiliated organisation; no external research grants were received. Total OpenRouter API cost across the entire project — generation pilots, audit cycles, and the final 16-config evaluation — was **\$783**, paid out-of-pocket.

A.2 Composition

What do the instances represent? Each instance is a multiple-choice benchmark question with: question text, 4 options with marked-correct answer, supporting fact(s) with source URL and tier label, generator model identifier, generation strategy, audit-agent verdicts (per-agent pass/warn/fail signals), and (for the gold-sheet subset) human-reviewer ratings on eight rubrics.

How many instances are there? **3,266** questions in `release_v1.2`. By strategy: fact-to-question 1,909, distractor mining 405, template 389, scenario 319, comparative 244. By difficulty (post-relabel): L1 694, L2 894, L3 678, L4 1,001. By domain: wine regions 1,108, grape varieties 766, producers 515, viticulture 502, wine business 250, winemaking 187. The corpus is built from 38,104 atomic facts across 35 scrapers and 4,295 individually-tracked source records spanning 56 unique top-level domains.

What data does each instance consist of? Structured Postgres records mirroring the schema documented in Section 3 and Appendix B.

Recommended data splits. The dataset ships as a single test split (this is an evaluation-only benchmark). Two analytic partitions are documented and recommended for use: (i) the *closed-book* slice (1,601 questions tagged `closed_book_solvable`) and the *contextual* slice (1,665 remaining questions), partitioned by audit agent B2; and (ii) per-generator held-out splits for self-preference analysis (Section 5.3), defined dynamically by the `generator_family` column.

390 **Errors, noise, redundancies?** Documented in Section 4. Per-agent fail rates and κ versus the
391 human gold sheet are released alongside the corpus. The 1,601 B2-flagged questions are kept with
392 disclosure because LLM-judge calibration with humans on closed-book solvability is unreliable
393 ($\kappa \approx 0.007$).

394 **Self-contained, or relies on external resources?** Self-contained: question text, supporting facts,
395 and source URLs are all included. Source URLs may rot over time; we mitigate by archiving Tier-1
396 source pages and recording the access timestamp.

397 A.3 Collection process

398 **How was the data acquired?** By 35 provenance-verified web scrapers (released in
399 `src/scrapers/`). Sources span government registries (INAO, TTB, OIV), inter-governmental
400 bodies, university research groups (notably UC Davis, USDA Extension), Wikipedia, Wikidata, peer-
401 reviewed journals (*OENO One*, *Vitis*, *AJEV*), and curated open datasets. Scraping was rate-limited
402 and used the disclosed user agent `OenoBench-Research/1.0` (academic wine benchmark).

403 **Time frame.** Fact extraction: 2026-03 through 2026-04. Question generation: 2026-04 through
404 2026-05-03. Evaluation: 2026-05-03.

405 **Ethical review.** The work is non-IRB-eligible at the authors’ institution because (a) no human
406 subjects were enrolled in research, (b) gold-sheet ratings were produced by three WSET-certified
407 reviewers, one of whom is among the authors, and (c) all source data is public-record information
408 about commercial wine entities. Human review of LLM-generated questions is treated as data-quality
409 assurance, not as human-subjects research.

410 A.4 Preprocessing / cleaning / labelling

411 Atomic-fact extraction with entity tagging; per-source paraphrase to avoid verbatim copy; near-
412 duplicate suppression at cosine ≥ 0.92 ; nine-agent automated audit; human gold-sheet rating on a
413 stratified 50-question subset per release, produced independently by three WSET-certified reviewers
414 (Diploma, Level 3, Level 2) with the senior domain-expert reviewer resolving disagreements.

415 A.5 Uses

416 **Recommended uses.** Evaluation of LLM factual knowledge in the wine domain; ablation studies
417 of bias-mitigation techniques in benchmark construction; study of self-preference effects in LLM-
418 as-judge pipelines; calibration of automated multi-agent QA frameworks against expert human
419 review.

420 **Discouraged uses.** Direct deployment as a wine-recommendation or purchasing system without
421 separate alignment, age-gating, and jurisdictional advertising review; fine-tuning on the benchmark
422 itself for ranking purposes (would invalidate the benchmark for the trained model). The scrape-to-
423 audit pipeline should not be retargeted to high-stakes domains (medicine, law) without comparable
424 expert human review of the retargeted source taxonomy.

425 **Alcohol context.** Wine is an alcoholic beverage. OenoBench is a knowledge-evaluation tool, not a
426 consumer-facing product. A high OenoBench score reflects factual / reasoning capability about wine;
427 it does not endorse downstream deployment as a drinking-recommendation system, which would
428 require separate alignment for health information, age-gating, and jurisdictional advertising rules.

429 A.6 Distribution

430 **Will the dataset be distributed?** Yes, under **CC-BY-SA-4.0**. The dataset, the Croissant manifest
431 [1], and the construction code are hosted at anonymous-review URLs disclosed in Appendix E; the
432 canonical post-acceptance HuggingFace dataset URL and GitHub code URL will be added to the
433 camera-ready version.

434 Citation.

```

435 @inproceedings{anonymous2026oenobench,
436   title = {OenoBench: A Wine-Domain Benchmark for Knowledge-Grounded
437           Evaluation of Large Language Models},
438   author = {Anonymous Author(s)},
439   booktitle = {Advances in Neural Information Processing Systems
440               (NeurIPS), Track on Evaluations and Datasets},
441   year = {2026},
442   note = {Under double-blind review; author identities withheld.}
443 }

```

444 A.7 Maintenance

445 **Maintainer.** Anonymous Author(s) during review; canonical maintainer contact will be disclosed
 446 in the camera-ready version.

447 **Update cadence.** Wine regulation changes on a multi-year cadence; we plan annual re-extraction.
 448 Scraper code is released so the community can re-extract independently. Corpus versioning follows
 449 `release_vX.Y` (major release / minor patch) with all prior versions retained on HuggingFace.

450 **Errata.** A public errata log will be maintained alongside the HuggingFace release; corrections to
 451 questions, facts, audit verdicts, or difficulty labels will be versioned and dated. Reports of factual
 452 errors are accepted via GitHub issues and will be evaluated against the source URL.

453 B Additional Methodology Detail

454 This appendix expands the dataset-construction and audit pipelines: the atomic-fact extractor (B.1),
 455 the Wikidata SPARQL choice (B.2), question-generation prompts (B.3), the closed-book gate (B.4),
 456 and audit architecture, agent specifications, and release-cycle results (B.5–B.7), with the calibration
 457 philosophy that informs them (B.8).

458 B.1 Atomic fact extraction pipeline




Figure 5: Atomic-fact extraction pipeline. Inputs are heterogeneous (HTML pages, RDF graphs, SPARQL responses, CSV dumps); outputs are uniform single-assertion sentences with entity tags, domain label, source URL and tier.

459 The atomic-fact extractor (`src/scrapers/_fact_processing.py`) applies five sequential filters
 460 to each candidate sentence:

- 461 1. **Sentence split.** NLTK Punkt tokeniser; further split at explicit conjunctions (“and”, “but”,
 462 “while”) when both clauses have an independent verb. Maximum atomic length 30 words;
 463 reject longer.
- 464 2. **Reference resolution.** Pronouns and demonstratives are replaced with their entity referent
 465 from article context (paragraph-level coreference using a rule-based resolver tuned on
 466 Wikipedia leads). Sentences with unresolved *it/this/that* are dropped.
- 467 3. **Domain classification.** Lexicon-based + rule-based classifier maps each fact to one of the
 468 six domain pillars; ambiguous facts (no anchor term in any domain lexicon) are dropped.
- 469 4. **Length and predicate validation.** Reject facts <5 words, >50 words, missing a verb, or
 470 carrying dangling references / fragments. The 5–30 word band is the strict filter; the 31–50
 471 band is kept with reduced confidence (910 facts).

472 5. **On-topic filter.** Region-specific keyword sets prevent cross-contamination (e.g. Austrian
473 content in a Bordeaux scraper). Fail $\Rightarrow$ drop.

474 A 2026-04 audit found that 19 of an original 35 scrapers contained hard-coded LLM-generated facts
475 disguised as scraped data. The fix required rebuilding all 19 against live sources; the lesson preserved
476 for future retargeting is that *provenance auditing must run on the scrapers themselves*, not only on
477 the collected facts.

478 B.2 Wikidata SPARQL: P17 vs P131*

479 A common bug in early scrapers used the transitive administrative-parent property (P131*) to bind
480 questions to a country. This introduced severe cross-region contamination, e.g. Austrian appellations
481 appearing in a Bordeaux scraper because the SPARQL graph traversal walked through shared ancestor
482 entities. Switching to the direct country relation (P17) eliminated cross-country leakage at the cost
483 of $\sim 3\%$ lost coverage on borderline-region entities. The released SPARQL templates use P17
484 exclusively and we recommend the same discipline for any pipeline retargeted to a domain with
485 hierarchical political subdivisions.

486 B.3 Question-generation prompts

487 The five strategies share a common JSON output schema enforced via Pydantic with a 3-tier extraction
488 fallback (markdown-fenced, raw JSON, prefix-pattern). Per-strategy prompt summaries:

- 489 • **Fact-to-question:** "Given a single atomic wine fact, write a 4-option multiple-choice
490 question whose unique correct answer is grounded in the fact. Paraphrase the fact; do not
491 copy contiguous spans of >5 words. Distractors must be confusable same-category entities,
492 not random alternatives."
- 493 • **Comparative:** given two related facts (matched on entity type, country/sub-domain), "ask
494 which differs in the named attribute or what both share. Avoid iconic-vs-non-iconic pairings."
- 495 • **Scenario synthesis:** given a cluster of 3–5 facts, "write a brief domain-appropriate pro-
496 fessional scenario (sommelier service / winemaker decision / viticulturist call / busi-
497 ness/regulatory choice) with a unique correct answer that requires *all* the cluster facts
498 to derive."
- 499 • **Distractor mining:** given a fact + 3 confusable entities, "write a question whose 3 wrong
500 options are these entities, ranked by ascending plausibility."
- 501 • **Template:** 45 deterministic parameter substitution templates; no LLM call. A subset is
502 paraphrased post-hoc by Gemini Pro to defeat the A4 stylometric fingerprint detector.

503 Full prompt texts are in `src/generators/_prompts.py`.

504 B.4 Closed-book gate specification

505 The closed-book pre-screen (`src/generators/_closed_book_gate.py`) runs a per-difficulty
506 model:

- 507 • L1: Claude Haiku 4.5 reads the question with no source fact. Correct answer $\Rightarrow$ tag
508 `closed_book_solvable`.
- 509 • L2: Claude Sonnet 4.6 same protocol.
- 510 • L3: Claude Opus 4.7 same protocol; the corpus quota cap fires here.
- 511 • L4: gate is skipped by protocol; expert items are expected to be answerable only with the
512 source fact.

513 The gate's role in generation is to enforce a corpus-level cap on closed-book solvable items; early
514 pilots used a 25% cap, but it was **raised to 50%** after the gold-sheet calibration revealed that the
515 LLM panel over-attributes closed-book solvability relative to humans (Section 4.2), so the higher
516 cap retains the items that are in fact discriminating in evaluation. The audit's B2 agent later re-runs
517 the check at corpus scale on a tri-judge panel. Failures are not dropped; they are reserved into the
518 closed-book slice and used for the analytic partition in Section 5.5.




Figure 6: Multi-agent audit architecture. Candidate questions are fanned out to four teams; each team independently emits per-question verdicts. Critical FAIL signals route the question to the drop pool; C4 difficulty disagreements route to relabel; otherwise the question is kept in release_v1.2. Every agent is calibrated against the human gold sheet via Cohen’s κ .

519 B.5 Audit architecture

520 B.6 Audit agent specifications

521 Each agent is a versioned Python module under `src/qa/agents/`. Agents emit (`question_id`,
522 `agent_id`, `version`, `severity`, `payload_json`) tuples to the `audit_findings` table;
523 `severity` is one of {PASS, WARN, FAIL}. Versioning is strict: the (`run_id`, `question_id`,
524 `agent_id`, `version`) tuple is unique, so re-running an agent at a new version produces additive
525 findings rather than overwriting. We document each agent’s threshold and gold-sheet κ here.

Table 5: Audit-agent thresholds and gold-sheet calibration.

Agent	Signal	Threshold	κ vs. gold
A1 LexicalHygiene	vague/marketing regex hits	≥ 1 hit $\Rightarrow$ WARN; $\geq 3 \Rightarrow$ FAIL	0.71
A2 BiasStats	χ^2 position; M–W length	$p < 0.01 \Rightarrow$ FAIL	0.65
A3 FactEcho	LCS / fact length	$\geq 0.45 \Rightarrow$ WARN; $\geq 0.65 \Rightarrow$ FAIL	0.83
A4 TemplateFingerprint	POS-bigram logreg held-out AUC	> 0.84 corpus-level $\Rightarrow$ WARN	0.62
B1 TriJudgeAnswer	majority disagrees with key	2/3 disagree $\Rightarrow$ FAIL	0.78
B2 ClosedBookSolv.	majority solves no-source	2/3 solve $\Rightarrow$ WARN; 3/3 $\Rightarrow$ FAIL	0.007
C2 CategoryLeak	wine-type mismatch in distractors	any $\Rightarrow$ FAIL	0.91
C4 DifficultyAudit	$ \Delta $ assigned vs. rated	≥ 1 WARN; ≥ 2 FAIL	0.74
B3 UbiquityRisk	ubiquitous-grape stem $\times$ region answer	both $\Rightarrow$ WARN	0.69

526 The B2 row is the calibration finding discussed in Section 4.2: the LLM panel and the human
527 reviewer disagree on what counts as “world-knowledge solvable”, because frontier LLMs have
528 absorbed substantial wine knowledge during pre-training. We retain the agent as a measurement of
529 the LLM panel’s prior, and report the contextual / closed-book partition (Section 5.5) that turns the
530 disagreement into a useful diagnostic.

531 B.7 Release-cycle results

532 Table 6 summarises the verdicts produced when the 9-agent audit was applied to release_v1.1
533 (3,670 candidate questions); these are the numbers cited in Section 4.3.

534 Table 7 shows the corpus-level effect of the C4 difficulty re-label, which shifted the corpus from 14%
535 L3+L4 items to 51% L3+L4.

Table 6: Audit findings on `release_v1.1` and the resulting drop / re-label policy.

Audit signal	Triggered FAIL	Action
A1 LexicalHygiene (vague phrasing)	60	drop
A3 FactEcho (verbatim copy)	63	drop
B1 TriJudgeAnswer (wrong answer key)	47	drop
C2 CategoryLeak (wine-category leak)	9	drop
B3 UbiquityRisk (grape $\times$ region ambiguity)	183	drop
Distinct dropped (with overlap)	341	
C4 DifficultyAudit (relabel from 1,252 LLM hits + 7 human)	1,259	re-label
B2 ClosedBookSolvability (kept w/ disclosure)	1,601	disclose

Table 7: Difficulty distribution before and after C4 re-label.

	Generator-assigned	Post-relabel	Δ
L1 (entry)	1,261 (38%)	694 (21%)	−567
L2 (intermediate)	1,559 (47%)	894 (27%)	−665
L3 (advanced)	218 (7%)	678 (21%)	+460
L4 (expert)	291 (9%)	1,001 (31%)	+710

536 B.8 Calibration philosophy

537 We emphasise the design choice underlying these numbers. A naive “automated audit” interpretation
538 would have dropped all 1,601 B2-flagged questions, which would have removed nearly half the
539 corpus and biased the remainder toward whatever idiosyncratic items the LLM panel happened not
540 to know. By calibrating B2 against humans and treating it as an evaluator-side signal, we preserve
541 dataset value while giving downstream users a principled way to interpret leakage. The audit’s role,
542 on this reading, is not to be the final arbiter of quality, but to surface signals *that humans validate* as
543 quality-relevant, and to disclose clearly the signals where humans and LLMs disagree.

544 C Full Evaluation Tables

545 This appendix gives the per-config tables and supporting figures behind each Section 5 subsection:
546 reasoning-mode lift (C.1), self-preference bootstrap CIs (C.2), the closed-book vs. source-grounded
547 contrast (C.3), and per-domain (C.4), per-strategy (C.5), and per-difficulty (C.6) breakdowns of the
548 overall ranking. The cost ledger that supports Section 5.4 is in Appendix E.




Figure 7: Reasoning-mode vs. standard accuracy for four families; δ in pp with 95% bootstrap CI (1,000 resamples). Only DeepSeek’s CI excludes zero.

Table 8: Reasoning lift per difficulty tier (δ = thinking – standard, pp).

Pair	L1	L2	L3	L4
Claude Opus 4.7 (think vs std)	+0.3	−0.2	+0.3	−0.5
Gemini 2.5 Pro (think vs std)	+0.5	+1.0	+1.0	+0.9
OpenAI o3 vs GPT-5	+0.1	−1.0	+1.9	+2.0
DeepSeek R1 vs V3	+1.7	+8.3	+5.7	+9.7

Table 9: Per-config Self-Preference Score with 95% bootstrap CI (1,000 resamples). n columns are the question counts in each pool: own = questions whose generator family matches the config’s family; other = questions whose generator is a different non-template LLM. Templates ($n=389$) are excluded from both pools, so $\text{own} + \text{other} = 2,877$ across every row. DeepSeek and Mistral have no own-family questions in release_v1.2 and are omitted.

Config	own n	other n	own (%)	other (%)	δ (pp)	95% CI
claude-haiku-4.5	619	2,258	58.8	48.8	+10.0	[+5.7, +14.5]
claude-opus-4.7	619	2,258	86.6	77.5	+9.1	[+6.1, +12.1]
claude-opus-4.7-thinking	619	2,258	86.9	77.2	+9.7	[+6.5, +13.1]
qwen-2.5-7b	667	2,210	60.9	51.9	+9.0	[+4.5, +13.2]
qwen-2.5-72b	667	2,210	66.1	64.1	+2.0	[−2.1, +6.0]
llama-3.3-70b	629	2,248	65.5	63.6	+1.9	[−2.3, +6.1]
gpt-5-mini	542	2,335	77.9	76.1	+1.7	[−2.4, +5.4]
o3	542	2,335	81.9	82.1	−0.1	[−3.9, +3.3]
gpt-5	542	2,335	79.7	81.4	−1.7	[−5.4, +2.1]
llama-3.1-8b	629	2,248	54.2	57.2	−3.0	[−7.3, +1.4]
gemini-2.5-pro	420	2,457	75.0	81.1	−6.1	[−10.7, −1.8]
gemini-2.5-pro-thinking	420	2,457	74.0	82.0	−8.0	[−12.3, −3.8]
gemini-2.5-flash	420	2,457	64.3	74.7	−10.4	[−15.3, −5.6]

Table 10: Per-tier Self-Preference Score δ (pp). The overall SPS in Table 9 is decomposed by the question’s post-relabel difficulty tier; each cell is own-Acc minus other-Acc on the questions of that tier. **Bold** marks cells whose 95% bootstrap CI excludes zero.

Config	Family	L1 δ	L2 δ	L3 δ	L4 δ
claude-haiku-4.5	anthropic	+3.1	+ 10.2	+2.0	+8.2
claude-opus-4.7	anthropic	+0.2	+ 8.5	+4.1	+ 9.5
claude-opus-4.7-thinking	anthropic	−0.6	+ 10.2	+3.4	+ 10.9
qwen-2.5-7b	qwen	+ 12.3	+6.6	+7.2	+ 10.4
qwen-2.5-72b	qwen	+2.4	−4.3	+0.0	+ 7.7
llama-3.3-70b	meta	−0.8	−3.8	−5.5	+6.8
gpt-5-mini	openai	+ 3.1	+ 10.7	+4.4	−0.7
o3	openai	+ 2.6	+ 6.9	+4.0	−3.2
gpt-5	openai	+ 2.9	+ 7.2	+0.8	−6.9
llama-3.1-8b	meta	−1.3	− 14.6	−6.9	+1.9
gemini-2.5-pro	google	−1.0	+4.5	+0.7	− 10.3
gemini-2.5-pro-thinking	google	+ 1.3	−0.3	+5.8	− 14.1
gemini-2.5-flash	google	−0.2	−6.3	−0.8	− 10.6

Per-tier reading. The Anthropic positive SPS is *not* concentrated at any single difficulty tier — it spikes at both L2 (+8.5 to +10.2 across the three Claude configs, all CIs above zero) and L4 (+9.5 to +10.9, two of three CIs above zero), while L1 (near ceiling) and L3 are flat. A pure question-difficulty asymmetry would predict the gap to widen monotonically as accuracy falls; instead the gap appears at *both* ends of the non-trivial range, which is more consistent with a residual stylistic fingerprint than with Anthropic having generated easier questions. Google’s negative SPS is overwhelmingly an L4 phenomenon (Gemini Pro / Pro thinking / Flash all −10 to −14 pp at L4, with CIs excluding zero; L1–L3 within ± 6 pp), consistent with Gemini-authored hardest-tier questions being uniformly hard rather than the model being penalised on its own style. OpenAI’s near-zero overall SPS is the result of positive L1/L2 and negative L4 cells averaging out, not stylistic neutrality.

C.3 Closed-book vs. source-grounded performance




Figure 8: Per-config accuracy on B2-flagged closed-book solvable questions ($n=1,601$) vs. contextual / source-grounded questions ($n=1,665$); every config gains +26.6 to +39.6 pp on the closed-book slice.

Table 11: Closed-book vs. contextual accuracy gap (pp) within each difficulty tier. A positive gap means the model scores higher on B2-flagged closed-book solvable items than on contextual items at the same tier. Per-tier B2-flagged / contextual question counts: L1 680 / 13, L2 333 / 561, L3 472 / 206, L4 116 / 885. The Mean row reports 95% bootstrap CIs (1,000 resamples; question-level resampling within each pool).

Config	L1 gap	L2 gap	L3 gap	L4 gap
claude-opus-4.7	+5.8	+30.0	+35.3	+31.0
claude-haiku-4.5	+13.0	+33.5	+22.6	+25.4
gpt-5	−1.8	+23.2	+21.0	+28.0
gpt-5-mini	−2.1	+32.8	+24.0	+26.1
gemini-2.5-pro	−1.5	+28.4	+27.8	+27.8
gemini-2.5-flash	+11.0	+33.9	+27.4	+28.5
llama-3.3-70b	+7.4	+37.4	+23.7	+19.1
llama-3.1-8b	−15.4	+29.0	+17.2	+11.6
deepseek-v3	+3.6	+41.0	+26.6	+31.4
qwen-2.5-72b	−6.9	+38.0	+25.0	+31.5
qwen-2.5-7b	−2.7	+30.9	+14.4	+21.3
mistral-large-2411	+2.1	+34.4	+28.6	+23.7
o3	−1.6	+24.8	+22.4	+29.6
gemini-2.5-pro-thinking	−0.9	+27.2	+25.1	+26.8
deepseek-r1	−2.5	+33.5	+28.5	+28.2
claude-opus-4.7-thinking	+13.9	+30.7	+32.3	+32.6
Mean (16 configs)	+1.3	+31.8	+25.1	+26.4
95% bootstrap CI	[−4.6, +8.6]	[+28.8, +34.8]	[+20.8, +29.1]	[+22.5, +30.0]

562 **Within-tier reading.** The +32.6 pp aggregate gap is preserved at L2–L4 (+31.8, +25.1, +26.4 pp;
563 every CI above zero) and effectively vanishes only at L1, where the contextual pool collapses to $n=13$
564 questions and ceiling effects dominate. The L4 result is the most informative: the contextual pool is
565 large ($n=885$) and the gap is still +26.4 pp with a tight CI, so the closed-book vs. source-grounded
566 contrast is not an artefact of difficulty composition — it is a stable property of the slice across
567 non-trivial tiers, consistent with B2 identifying questions recoverable from pre-training rather than
568 questions that happen to be easier-stated.

569 C.4 Per-domain breakdown

Table 12: Per-config per-domain accuracy (%).

Config	Regions	Varieties	Producers	Viticulture	Winemaking	Business
claude-opus-4.7	83.2	80.2	86.8	77.9	86.6	64.2
claude-haiku-4.5	55.8	49.7	58.7	54.2	56.7	37.8
gpt-5	88.1	78.2	88.4	82.8	80.7	63.0
gpt-5-mini	82.1	74.4	83.5	77.5	77.5	66.3
gemini-2.5-pro	87.0	78.8	89.0	78.5	77.5	61.8
gemini-2.5-flash	76.9	72.9	82.3	73.6	77.5	60.6
llama-3.3-70b	68.9	62.9	77.2	65.7	70.6	50.4
llama-3.1-8b	62.2	55.6	69.5	61.7	59.4	47.6
deepseek-v3	74.4	68.5	73.2	70.8	69.0	51.2
qwen-2.5-72b	69.0	65.1	73.4	67.5	70.1	52.8
qwen-2.5-7b	59.1	55.6	59.4	57.4	57.8	44.7
mistral-large-2411	71.5	65.6	75.4	69.4	73.3	51.6
o3	88.4	81.3	90.9	80.3	76.5	65.4
gemini-2.5-pro-thinking	86.5	80.5	90.0	80.5	83.4	60.2
deepseek-r1	80.5	73.6	85.6	72.2	76.5	64.6
claude-opus-4.7-thinking	83.2	80.9	87.0	78.7	84.0	61.4
Mean (16 configs)	76.0	70.2	79.4	71.8	73.6	56.5




Figure 9: Per-domain accuracy as a heatmap. *Wine business* is the hardest domain across every tier (mean 56.5%) — both because business-domain facts are sparser in the training distribution and because the C4 difficulty re-label promoted business questions that are computationally harder (multi-fact reasoning over regulation).

570 C.5 Per-strategy breakdown

Table 13: Per-config per-strategy accuracy (%).

Config	FTQ	Scenario	Template	Comparative	Distractor
claude-opus-4.7	77.7	85.6	92.8	79.5	82.7
claude-haiku-4.5	47.5	62.7	70.7	53.3	56.8
gpt-5	79.9	80.9	95.4	79.5	87.9
gpt-5-mini	73.8	80.6	92.8	78.3	84.7
gemini-2.5-pro	77.4	87.5	93.1	82.4	86.4
gemini-2.5-flash	69.5	84.0	89.7	77.5	79.5
llama-3.3-70b	59.5	80.6	89.7	64.8	71.9
llama-3.1-8b	52.2	70.5	89.7	60.2	63.7
deepseek-v3	64.5	80.3	86.4	71.7	73.3
qwen-2.5-72b	59.6	79.3	88.7	69.3	73.6
qwen-2.5-7b	49.1	72.1	78.9	57.0	61.0
mistral-large-2411	62.7	80.6	87.4	67.6	73.3
o3	79.9	84.3	94.9	84.8	88.4
gemini-2.5-pro-thinking	78.5	88.4	95.6	84.0	84.2
deepseek-r1	73.4	75.5	92.5	77.0	80.7
claude-opus-4.7-thinking	77.3	85.6	93.6	80.3	83.0
Mean (16 configs)	67.6	79.9	89.5	73.0	76.9




Figure 10: Per-strategy accuracy heatmap. Templates are the easiest strategy uniformly (89.5% mean), even for small models, because they are deterministic single-fact recall questions. The fact-to-question strategy is hardest because its distractors are sampled from confusable same-category entities; comparative questions add the constraint that the model must integrate two facts.

571 C.6 Per-difficulty breakdown

Table 14: Per-config per-difficulty accuracy (%).

Config	L1 (entry)	L2	L3	L4 (expert)
claude-opus-4.7	98.0	77.0	86.7	69.1
claude-haiku-4.5	74.3	48.3	59.9	38.8
gpt-5	98.3	82.4	87.5	69.2
gpt-5-mini	98.0	73.7	84.7	64.8
gemini-2.5-pro	98.6	79.2	87.3	68.5
gemini-2.5-flash	95.4	70.9	82.2	60.1
llama-3.3-70b	91.9	60.6	74.8	50.3
llama-3.1-8b	84.8	55.4	65.3	45.0
deepseek-v3	95.8	65.9	76.3	52.4
qwen-2.5-72b	93.2	60.5	73.7	51.4
qwen-2.5-7b	82.0	50.6	60.0	43.3
mistral-large-2411	94.4	63.1	73.3	54.0
o3	98.4	81.4	89.4	71.2
gemini-2.5-pro-thinking	99.1	80.2	88.3	69.4
deepseek-r1	97.5	74.2	82.0	62.1
claude-opus-4.7-thinking	98.3	76.8	87.0	68.6
Mean (16 configs)	93.6	68.8	78.7	58.7




Figure 11: Difficulty calibration: distribution of accuracy across the 16 configs at each level. The medians sit at L1 96%, L2 70%, L3 81%, L4 60%.

On the L2 / L3 / L4 means in Table 14. Two features of the column averages deserve comment. First, the L2 mean (68.8%) is *below* the L3 mean (78.7%) — an apparent inversion of the difficulty ordering. This is an artefact of the C4 difficulty re-label (Section 4): many questions originally labelled L2 by their generator were promoted to L3, leaving behind a residual L2 set that is concentrated in the *wine_business* pillar (the hardest domain by a wide margin, Table 12). The L3 set, in contrast, was enriched with cleanly-stated multi-fact items from the higher-quality scrapers. Second, the L3 → L4 drop is real and large (78.7% → 58.7%, a 20 pp fall), reflecting that L4 is the expert tier where the source fact is genuinely required to answer the question; that fall is also where the inter-config spread is widest (Section 5.1), making L4 the most discriminating slice for distinguishing models. Both effects are intended outcomes of the audit-driven difficulty re-label rather than miscalibration.

D Application Screenshots

D.1 Human-review web application

The human-review web application is the front-end through which WSET-Diploma reviewers rate stratified gold-sheet questions and the through which the agent-vs-human κ values reported in Section 4.2 are collected. It is a Flask application with two layers of authentication (outer HTTP Basic for shared-link access, inner per-reviewer session cookies) and writes ratings into a separate `human_reviews` Postgres table to keep gold-sheet labels independent of automated audit findings. Each reviewer is associated with a WSET certification level recorded at registration, and the app supports multi-reviewer inter-rater agreement (κ) computation out of the box.




Figure 12: Sign-in / register screen. Outer HTTP Basic Auth gates shared-link access; the inner form establishes a per-reviewer session for IRR attribution.




Figure 13: Reviewer dashboard. Active batches are listed; clicking a batch enters the rubric-scoring flow (eight-rubric grid with PASS/WARN/FAIL per rubric, inline reviewer-guide hints, and time-on-page tracking).

E Released Artifacts and Reproducibility

E.1 Released artifacts

For double-blind review, the canonical HuggingFace and GitHub URLs are withheld; the artifacts listed below are mirrored at anonymous review URLs and will be relinked to the canonical hosts in the camera-ready version.

- **Dataset (anonymous mirror).** [anonymous-dataset-url-withheld-for-review] — release_v1.2, 3,266 questions, Parquet test split, README.md datasheet, croissant.json manifest. License: CC-BY-SA-4.0. Versioning is permanent: prior releases remain accessible via revision tags.
- **Code (anonymous mirror).** [anonymous-code-url-withheld-for-review] — full pipeline: 35 scrapers, 5 question generators, 14 audit agents, the human review web app, and the monitoring dashboard. Apache 2.0.
- **Process log.** docs/PROCESS_LOG.md — chronological lab-notebook of every phase shipped, with sources, methodology, quality-control counts, and decision rationale, in the format prescribed for the methodology sections of this paper.
- **Audit reports.** docs/QUALITY_AUDIT_REPORT.md, docs/RELEASE_V1_2_AUDIT_ACTIONS.md, docs/GOLD_CALIBRATION_ANALYSIS.md, data/reports/zero_correct_97_audit.md.
- **Eval reports.** data/reports/eval_release_v1_2_full_cleaned.md (the 16-config × 3,266-Q run reported in this paper).

E.2 Reproducibility scripts

The construction pipeline is driven by a small number of shell scripts in scripts/:

- run_all_scrapers.sh — orchestrates the 35 scrapers end-to-end. Each scraper writes a timestamped log to data/logs/.
- run_release_v1_build.sh — full generation pipeline invocation; idempotent and resume-safe via stamped tags.
- export_release_v1_2_to_parquet.py — produces the HuggingFace Parquet directly from Postgres; the released file is the byte-identical output of this script.
- build_smart_review_sheet.py — generates a 50-question stratified gold-sheet review batch (random + critical-FAIL + borderline-WARN strata).
- tag_audit_actions.py — applies the audit drop / relabel policy to a tagged corpus; this is the script that produced the release_v1.2 state from release_v1.1.

E.3 Database schema

The released schema is initialised by the migrations under `config/postgres/` (six idempotent files, applied in order on container startup): `init.sql` (sources, facts, questions, question_facts, generation_metadata, evaluation_runs, evaluation_answers, human_reviews), `002_audit_schema.sql` (audit_runs, audit_findings, audit_gold_labels, audit_severity enum), `003_sample_schema.sql` (mirror schema for the Phase 5 sample-DB eval), `003_cb_reserve.sql` (closed-book reserve pool tagging), `004_eval_telemetry.sql` (provider, tokens, reasoning_config, latency columns on evaluation_answers), `005_or_cost_telemetry.sql` (OpenRouter authoritative cost columns).

The released schema is documented in code (`config/postgres/`); the load-bearing tables are `sources` (registry), `facts` (atomic facts grounded by source), `questions + question_facts` (corpus and fact-question links), `generation_metadata` (per-question generator and strategy), `audit_findings` (per-agent verdicts), and `human_reviews` (gold-sheet ratings).

E.4 Compute and cost

The full evaluation reported in Section 5 ran in **120 minutes 34 seconds** on a single VM driving 16 OpenRouter configurations in parallel. **Total LLM API spend across the entire project was \$783**, paid out-of-pocket to OpenRouter and covering all generation pilots, audit cycles, gold-sheet calibration runs, and the final 16-config evaluation. No GPU compute was used; all model inference was via OpenRouter (a unified API gateway). The release artefacts include all per-call telemetry (provider, tokens, reasoning configuration, p50/p95 latency, OR-authoritative cost) for the full evaluation run; the per-config cost ledger is reproduced inline as Table 15.

Table 15: Per-config cost ledger for the full 16-config evaluation (3,266 questions per config). Effective OpenRouter spend.

Config	Total cost	Correct ans.	Cost / correct
llama-3.1-8b	\$0.01	1,976	\$0.0000
llama-3.3-70b	\$0.04	2,190	\$0.0000
gemini-2.5-flash	\$0.12	2,454	\$0.0000
deepseek-v3	\$0.12	2,295	\$0.0001
qwen-2.5-7b	\$0.12	1,860	\$0.0001
qwen-2.5-72b	\$0.15	2,202	\$0.0001
claude-haiku-4.5	\$0.44	1,741	\$0.0003
mistral-large-2411	\$0.87	2,256	\$0.0004
gpt-5-mini	\$2.82	2,561	\$0.0011
claude-opus-4.7	\$3.35	2,647	\$0.0013
claude-opus-4.7-thinking	\$3.53	2,645	\$0.0013
deepseek-r1	\$10.54	2,517	\$0.0042
o3	\$11.80	2,729	\$0.0043
gemini-2.5-pro-thinking	\$13.06	2,698	\$0.0048
gpt-5	\$21.90	2,704	\$0.0081
gemini-2.5-pro	\$29.47	2,669	\$0.0110

E.5 Per-source licensing

Section 3.2 states that every fact in the corpus traces to a public-record source or to a source distributed under a licence that permits scraping and redistribution. We give the per-source breakdown in Table 16, grouping the 35 scrapers’ source families by tier and licence basis. Each row maps to one or more rows in the released `sources` table, where per-row `name`, `url`, `tier`, and `access_timestamp` are recorded for every fact (`facts.source_id` → `sources.id`). Wikipedia and Wikidata together account for 65.3% of facts and propagate their share-alike / public-domain dedications through the `release_v1.2 CC-BY-SA-4.0` licence. Government records (Tier 1, 19.6% of facts) are either explicit open-data licences (INAO Licence Ouverte 2.0, EU re-use Decision 2011/833/EU) or US-Government works in the public domain (17 U.S.C. §105: TTB, USDA Extension, UC IPM).

The released corpus, including this paper’s quoted statistics, is distributed under **CC-BY-SA 4.0**. The construction code is released under **Apache 2.0**. Verbatim source text is never stored: every fact is paraphrased atomically (Section 3.3) and a paraphrase guard at audit time (agent A3) re-checks

Table 16: Per-source licensing for OenoBench’s 35 scrapers, grouped by source family. “Share” is the percentage of the 38,104-fact corpus contributed; “Tier” is the source-of-authority label from Section 3.2. Counts are computed directly from the `facts` $\bowtie$ `sources` join in the released schema.

Source family	Tier	Licence / legal basis	Share
Wikipedia (MediaWiki API)	T2	CC-BY-SA 4.0	34.4%
Wikidata (SPARQL)	T1 / T2	CC0 1.0	30.3%
HuggingFace <code>spawn99/wine-reviews</code>	T2	HuggingFace dataset terms (open re-distribution)	8.5%
UC Davis (AVA Project, FPS, Wine Ontology)	T1	Open / CC-BY (UC Davis Library)	5.7%
UC IPM & USDA Extension (eXtension, PSU/OSU)	T1 / T2	Public domain — US Gov. work (17 USC §105)	4.7%
Kaggle: wine-reviews (Wine Enthusiast scrape)	T3	Kaggle dataset terms (research use)	3.8%
INAO (data.gouv.fr open-data)	T1	Licence Ouverte 2.0 (Etalab)	3.9%
Wine consortia + national bodies (CIVB, BIVB, CIVC, IT consortia, NZ Wine, Austrian Wine, Deutsches Weininstitut, Spanish DO)	T2	Industry-body publications, educational use	3.2%
OENO One	T2	CC-BY 4.0 (open-access journal)	2.0%
TTB (27 CFR Part 4, Established AVAs, Approved Grape Names)	T1	Public domain — US Gov. work	1.3%
National ministries + EU registries (MASAF, MAPA, IVV, IVDP, EUR-Lex, GIView)	T1 / T2	Government public records / EU Decision 2011/833/EU	1.3%
OIV publications (Code of Oenological Practices, statistics)	T1	Open access — intergov. organisation	0.2%
Vitis (Hochschule Geisenheim)	T2	Open access	0.3%
National reference databases (CA, UK, HR/SI, HU/GE, LB/IL, +regions)	T2	Reference / official-body content	0.3%
Kaggle: wine-quality (UCI ML Repository)	T2	CC-BY 4.0 (UCI)	0.2%

656 for verbatim leakage. We retain `accessed_date` and a per-source `content_date` for every fact, so
657 dataset users can audit freshness against original sources.

658 E.6 Test status

659 771 of 771 unit tests pass on `main` as of the camera-ready build. Tests cover the fact-processing
660 pipeline, the five generation strategies, all 14 audit agents, the orchestrator, the closed-book gate, the
661 eval harness (slot dispatch, cost computation, reasoning-mode prompt assembly), and the review-app
662 endpoints.

663 F Construction Timeline

664 The full chronological lab notebook of construction is in `docs/PROCESS_LOG.md`. We summarise
665 the major dated milestones here for paper-traceable provenance.

Table 17: Construction-and-evaluation timeline (selected milestones).

Phase	Dates	Milestone
0	2026-02-13	Repository initialised: Docker stack (Postgres / Elasticsearch / Neo4j / Redis), schema, and the first four scrapers (Wikidata, HuggingFace, Kaggle, UC Davis) committed.
1	2026-02-13 – 04-01	35 scrapers built incrementally (INAO, TTB, EU/OIV, Wikipedia, Italian / Spanish / Portuguese / German / Austrian / Greek registries, Bordeaux / Burgundy / Champagne / Italian consortia, USDA Extension, OENO One / Vitis / AJEV academic crawlers, country enrichment passes); first ~16,702-fact working corpus assembled by end of March.
1.1	2026-04-07	<i>Provenance audit reveals 19 scrapers contained hardcoded LLM-generated facts</i> ; corpus paused.
1.2	2026-04-08 – 04-11	Shared infrastructure (<code>_fact_processing.py</code> , <code>_web_helpers.py</code> , <code>_wiki_helpers.py</code>) shipped; Phase 1 rebuild fixes 8 quality-issue scrapers; Phase 2 rebuild replaces all 17 hardcoded scrapers with genuine Wikipedia + SPARQL + official sources. <i>38,104 fact corpus restored, all provenance-traceable</i> .
2	2026-04-12 – 04-22	Five generation strategies built and iteratively tuned; multi-model orchestrator with quota management; closed-book gate v2 wired into all five generators.
2c	2026-04-18 – 05-02	Nine-agent audit framework shipped across 4 teams; sixteen audit pilots (v1–v16) tune thresholds and speedups; 472/472 → 771/771 tests pass.
2g	2026-04-28 – 05-02	Ten generation-pipeline speedup levers shipped (LLM cache, in-process dispatch, Thread-Pool dispatch, substantiveness filter, circuit breaker, tier-aware gate). v9 build wall: 18 min vs. v8’s 11.4 h.
2j	2026-05-03	<code>release_v1.1</code> (3,670 candidate questions) assembled from staged generation runs and the curated sample-DB. Audit cycle drops 341 questions on critical FAILs and re-labels difficulty for 1,259 (51% L3+L4 vs. prior 14%). Cycle cost \$76 / 5h 22m.
4	2026-05-03	Human-review web app shipped on port 5556 with separate <code>human_reviews</code> schema and multi-reviewer IRR support.
5	2026-05-03 – 05-04	Full evaluation: 16 configs × 3,329 questions in 120m 34s for \$98.33; post-eval audit of 97 zero-correct items removes 54 defects + 9 borderline-review drops; <code>release_v1.2</code> locked at 3,266 questions.
HF	2026-05-04	Dataset card + Croissant manifest + Parquet test split published; this paper drafted.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: Section 1 lists four contributions, and each is supported by a numbered section: dataset construction (Section 3), automated audit (Section 4), bias-aware evaluation (Section 5). The abstract’s headline numbers (3,266 questions, 16-config eval, +32.6 pp closed-book vs. contextual gap) are reproduced with full tables in Appendix C.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Section 7 discusses three limitations explicitly: snapshot vs. moving target, B2 audit calibration ceiling, and residual self-preference.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[N/A\]](#)

Justification: This is a dataset-and-benchmark paper; no formal theoretical results are claimed.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper?

Answer: [\[Yes\]](#)

Justification: Appendix E lists the construction scripts, the dataset URL, the model slate, the per-call telemetry available alongside the released corpus (including run id, wall time, OR-authoritative cost, and exact config strings), and the deterministic build pipeline. Section 5 specifies the 16-config slate, the single-letter (A–D) output protocol, and the bootstrap procedure.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [\[Yes\]](#)

Justification: Code (Apache 2.0) and data (CC-BY-SA-4.0) are released at the anonymous review mirrors documented in Appendix E; canonical GitHub and HuggingFace URLs are withheld for double-blind review and will be added to the camera-ready version. Exact reproduction commands are documented in Appendix E.

6. Experimental setting/details

Question: Does the paper specify all the training and test details necessary to understand the results?

Answer: [\[Yes\]](#)

Justification: Section 5 specifies the 16-config slate, single-letter output protocol (A–D, max_tokens=5, five-stop fallback). Appendix C gives full per-config × per-domain × per-strategy × per-difficulty tables. Appendix B gives audit-agent thresholds and prompt summaries.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [\[Yes\]](#)

Justification: Reasoning lift, self-preference scores, and closed-book vs. contextual deltas are all reported with 95% bootstrap confidence intervals (1,000 resamples). Section 5 states the procedure; Appendix C (Tables 8, 9, 10, 11) reports per-config CIs.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources needed to reproduce the experiments?

Answer: [\[Yes\]](#)

Justification: Appendix E reports the full 16-config evaluation wall time (120 min 34 s), total LLM calls (52,256), and effective evaluation cost (\$98.33), and the deployment topology (single VM driving 16 OpenRouter configurations in parallel; no GPU compute). The total OpenRouter spend across the entire project — generation pilots, audit cycles, gold-sheet calibration, and the final evaluation — was **\$783**, paid out-of-pocket and itemised in Appendix E and the Datasheet (Appendix A, Motivation).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics?

Answer: [\[Yes\]](#)

Justification: All sources used are public-record; no human subjects research was conducted (gold-sheet ratings were produced by three WSET-certified reviewers, one of whom is among the authors, treated as data-quality assurance); responsible-use guidance is given in the Datasheet (Appendix A, “Uses”) and in Section 7.

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [\[Yes\]](#)

Justification: The Datasheet (Appendix A) covers alcohol-context and discouraged uses; source-licensing and the non-fabrication guarantee are documented in Section 3.2 and Table 16 (Appendix E); the corpus contains only public-record information about commercial wine entities (no private individuals; Datasheet “Collection process”). Section 6 discusses positive applications including specialist-model fine-tuning and pipeline retargeting to other expert domains.

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse?

Answer: [Yes]

Justification: The dataset is text-only multiple-choice questions with public-record entity names; no scraped images or personal-data risks. We document discouraged uses in the Datasheet (Appendix A, “Uses”) including the recommendation against retargeting the pipeline to high-stakes domains without expert review.

12. Licenses for existing assets

Question: Are the creators or original owners of assets used in the paper properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The full per-source licensing table (Table 16 in Appendix E) credits every source family with licence and legal basis (Wikipedia CC-BY-SA 4.0, Wikidata CC0, INAO Licence Ouverte 2.0, TTB / USDA / UC IPM public domain (US Government work, 17 USC §105), EU re-use Decision 2011/833/EU, OENO One CC-BY 4.0, etc.) and is also summarised in Section 3.2. Per-fact source URL, accessed-date, and content-date are released alongside the corpus.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: A full Datasheet is in Appendix A; a Croissant manifest [1] ships alongside the released dataset (anonymous mirror during review, canonical HuggingFace at camera-ready); all 14 audit agents (9 always-run, 5 escalation-gated; Section 4.1), 5 generation strategies, 5 generator-model families, and 35 scrapers are documented in source.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation?

Answer: [N/A]

Justification: No crowdsourcing or third-party human-subjects research. Gold-sheet ratings were produced by three WSET-certified reviewers (Diploma, Level 3, Level 2), one of whom is among the authors; the review web app screenshots in Appendix D document the rubric and instructions.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether IRB approvals were obtained?

Answer: [N/A]

Justification: As noted in Appendix A, the work involves no human subjects research distinct from data-quality assurance by three WSET-certified reviewers, one of whom is among the authors.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research?

793
794
795
796
797
798

Answer: [\[Yes\]](#)

Justification: LLMs are central to the methodology (multi-strategy question generation across five model families, nine-agent automated audit including a tri-judge panel, closed-book solvability pre-screen). Section 3 and Section 4 describe their usage; Section 3.2 and Appendix B document the explicit fact-grounding constraint that prevents LLMs from being used as the source of factual claims.